

Lecture 13

Name: Tyler LaBonte

Professor: Santosh S. Vempala

1 Weighted Majority

Recall that the perceptron algorithm has mistake bound

$$M \leq \frac{\|\mathbf{w}^*\|_2^2 \|\mathbf{x}^*\|_2^2}{\gamma^2} \propto \frac{1}{\gamma^2}. \quad (1)$$

for learning halfspaces. In a different setting, suppose we are given the predictions of n experts (*e.g.*, data sources) which we can use to make our decision. This information can still be encoded as a point in Euclidean space similarly to the perceptron setting, and we will indeed see a connection to halfspaces at the end of the lecture. The expert setting is more general, though, because we could imagine the experts as paths in a graph (exponential) or probability distributions (infinite). Additionally, we can use bandit feedback or full feedback.

A reasonable goal is to do as well as the best expert. We first try majority vote, *i.e.*, we predict 1 if at least half the experts predict 1, and we predict 0 otherwise. Say we make M mistakes and there is a perfect expert, and our learning strategy is to throw out the majority if they are wrong. Then $M \leq \log n$ since we can make at most $\log n$ mistakes before only the perfect expert remains.

This is obviously unrealistic, so let's say the best expert makes m mistakes and that we restart after we throw out all the experts. Then, for every $\log n$ mistakes of the algorithm, each expert makes at least 1 mistake. So, $M \leq m \log n$. This is still a problem if n is large.

We will employ the weighted majority (multiplicative weights) algorithm to do better:

1. Initialize $w_i \geq 0$ at $w_i = 1$ for all experts i .
2. Predict according to the weighted majority.
3. On a mistake, for each expert i who got it wrong, update $w_i \leftarrow w_i/2$.

We can see that $W = \sum_i w_i$ starts at 1, and on a mistake, it goes down to $3/4$ of its current size (because we halve the majority which is at most half the total weight). So after M mistakes, $W \leq n \left(\frac{3}{4}\right)^M$. And since the best expert makes at most m mistakes, $W \geq \left(\frac{1}{2}\right)^m$. So,

$$\left(\frac{1}{2}\right)^m \leq n \left(\frac{3}{4}\right)^M. \quad (2)$$

Taking logs,

$$-m \leq \log n - M \log \frac{4}{3} \quad (3)$$

$$\implies M \leq \frac{m + \log n}{\log 4/3} \quad (4)$$

$$\approx \frac{3}{2}(m + \log n) \quad (5)$$

We can do even better by randomizing. Suppose we take a parameter ϵ , predict according to expert i with probability w_i/W , and for each expert that makes an error, update $w_i \leftarrow (1 - \epsilon)w_i$. Notice that this update happens for each expert that gets it wrong every round, not just when our algorithm makes a mistake.

Define f_t to be the weighted fraction of experts that got it wrong at time t . Then $\mathbb{E}[M] = \sum_{t=1}^T f_t$. The total weight at time T is

$$W = n \prod_{t=1}^T (1 - \epsilon f_t) \quad (6)$$

and the weight of the best expert i is

$$w_i \geq (1 - \epsilon)^m. \quad (7)$$

Taking logs and using the fact that $1 + x \leq e^x$,

$$m(1 - \epsilon) \leq \ln n + \sum_{t=1}^T (1 - \epsilon f_t) \quad (8)$$

$$\leq \ln n - \epsilon \sum_{t=1}^T f_t \quad (9)$$

$$= \ln n - \epsilon \mathbb{E}[M]. \quad (10)$$

So,

$$\mathbb{E}[M] \leq (1 + \epsilon)m + \frac{1}{\epsilon} \ln n. \quad (11)$$

We can differentiate to find the optimal ϵ is $\sqrt{\frac{\ln n}{m}}$. Thus,

$$\mathbb{E}[M] \leq m + 2\sqrt{m \ln n}. \quad (12)$$

We can analyze the average expected error to further understand this quantity. Using $m \leq T$,

$$\frac{\mathbb{E}[M]}{T} \leq \frac{m}{T} + \frac{2\sqrt{m \ln n}}{T} \leq \frac{m}{T} + 2\sqrt{\frac{\ln n}{T}}. \quad (13)$$

The second term goes to 0 as $T \rightarrow \infty$, so we converge to the best expert.

2 Winnow Algorithm

We can extend the weighted majority idea to our supervised binary classification setting with the winnow algorithm. Suppose $\mathbf{x} \in \{0, 1\}^n$ and $y \in \{0, 1\}$. Then,

1. Initialize $w_i \geq 0$ at $w_i = 1$ for all i .
2. Predict 1 if $\mathbf{w}^\top \mathbf{x} \geq n$ and 0 otherwise.
3. On a mistake:
 - a) If the true label is 1, then for all i with $x_i = 1$, update $w_i \leftarrow 2w_i$.
 - b) If the true label is 0, then for all i with $x_i = 1$, update $w_i \leftarrow w_i/2$.

Note that we do not update the w_i for which $x_i = 0$.

Let's start with an easy example. Suppose we apply the winnow algorithm to a disjunction (OR) or r out of n variables, called the relevant variables. Note that the weight of the relevant variables never goes down and is at most n (since then we will always predict correctly). So, the number of positive mistakes $M_+ \leq r \log n$, since at that point we will always predict correctly.

On every positive mistake, the total weight increases by at most n , and on every negative mistake, the total weight decreases by at least $n/2$. Since the total weight must remain positive we have $M_- \leq 2M_+$. So $M = M_- + M_+ \leq 3r \log n$.

Now, suppose we apply the winnow algorithm to a k out of r majority function, where there are r relevant variables and we need k of them for the label to be 1. If $k = 1$ is it a disjunction, if $k = r$ it is a conjunction, and if $k = r/2$ it is a simply majority. We will also add a small change in our update rule: replace 2 by $1 + \epsilon$ for a more general bound.

On a positive mistake, at least k variables get a $1 + \epsilon$ increase, and on a negative mistake, at most $k - 1$ variables get a $1 - \epsilon$ decrease. Since the weight of the relevant variables cannot exceed n , we obtain

$$kM_+ - (k - 1)M_- \leq r \log_{1+\epsilon} n. \quad (14)$$

And since the weight must be positive,

$$n + \epsilon n M_+ \geq \frac{\epsilon n}{1 + \epsilon} M_-. \quad (15)$$

Rearranging,

$$M_- \leq (1 + \epsilon)M_+ + \frac{1 + \epsilon}{\epsilon} n. \quad (16)$$

Plugging into the first equation,

$$(k - (1 + \epsilon)(k - 1))M_+ \leq r \log n + (k - 1)\frac{1 + \epsilon}{\epsilon} n. \quad (17)$$

We can again differentiate to find the optimal $\epsilon = \frac{1}{2(k-1)}$. So, $M = \mathcal{O}(kr \log n)$.

3 Winnow for Halfspaces

Previously we had the conditions $\sum_{i=1}^r x_i \geq 1$ and $\sum_{i=1}^r x_i \geq k$. It is actually simple to extend the winnow algorithm to general halfspaces $\sum_{i=1}^n w_i^* x_i \geq w_0^*$. We can assume $w_i^* \geq 0$ by negating the variable if necessary, and we can assume $w_i^* \in \mathbb{Z}$ by scaling (as long as they are rational to begin with). Furthermore, we can duplicate each variable $W = \sum_i w_i^*$ times. Now, we have a “learn w_0^* out of W ” problem – we just solved this!

Plugging into the bound we just found,

$$M = \mathcal{O}(w_0^* W \log(nW)) = \mathcal{O}(W^2 \log(nW)). \quad (18)$$

More generally,

$$M = \mathcal{O}\left(\frac{\|\mathbf{w}^*\|_1^2 \|\mathbf{x}\|_\infty^2 \log(n\|\mathbf{w}^*\|_1)}{\gamma^2}\right). \quad (19)$$

Interestingly, both winnow and perceptron have their own advantages. If \mathbf{w}^* is sparse, say $(1, 1, \dots, 1, 0, 0, \dots, 0) \in \mathbb{R}^n$, then winnow gives $M = \mathcal{O}\left(\frac{k^2 \log n}{\gamma^2}\right)$ and perceptron gives $M = \mathcal{O}\left(\frac{kn}{\gamma^2}\right)$. But, if \mathbf{w}^* is dense, say $(1/\sqrt{n}, 1/\sqrt{n}, \dots, 1/\sqrt{n})$, then winnow gives $M = \mathcal{O}\left(\frac{n \log n}{\gamma^2}\right)$ while perceptron gives $M = \mathcal{O}\left(\frac{1}{\gamma^2}\right)$.